


# Sequence modeling – recurrent networks


Biplab Banerjee

# Sequential learning problem - motivation



- ✓ Each input is of fixed size
- ✓ Inputs are independent

# Sequential learning problem




Auto-correct application in e-mail

- ✓ Input lengths can vary
- ✓ Inputs are no longer independent
- ✓ At each time stamp, the same operation is Carried out

Given previous **information**  
And current input, **predict**  
Next input

# Sequential learning problem



Parts of speech tagging

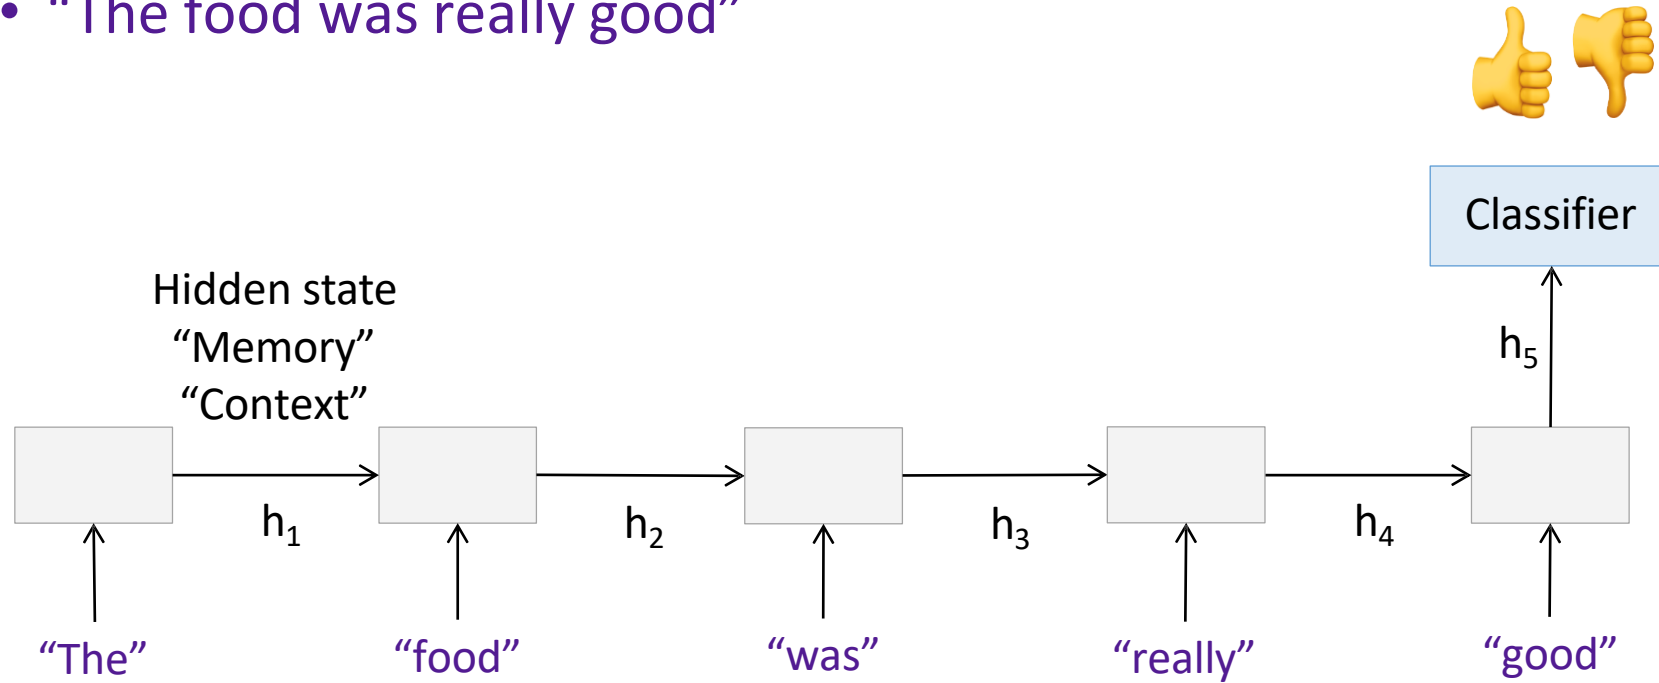
A set of previous time-stamps provide useful information  
About the next time-stamp

# Text classification

- **Sentiment classification:** classify a restaurant or movie or product review as positive or negative
  - “The food was really good”
  - “The vacuum cleaner broke within two weeks”
  - “The movie had slow parts, but overall was worth watching”
- What feature representation or predictor structure can we use for this problem?

# Sentiment classification

- “The food was really good”



Recurrent Neural Network (RNN)

# Machine translation

The screenshot shows the Google Translate web interface. At the top is the Google logo. Below it, the word "Translate" is in red. To the right, there is a link "Turn off instant translation" and a star icon. The main area has two language selection bars. The left bar shows "English", "Spanish", "French", and "Detect language" with a dropdown arrow. The right bar shows "English", "Spanish", "Arabic", and a dropdown arrow, followed by a blue "Translate" button. Below these bars, the text is displayed in two columns. The left column contains the original French text, and the right column contains the English translation. At the bottom left, there are icons for voice input, a keyboard, and a language dropdown, along with the text "693/5000".

Google

Translate

Turn off instant translation

English Spanish French Detect language

English Spanish Arabic Translate

Correspondances

La Nature est un temple où de vivants piliers  
Laissent parfois sortir de confuses paroles;  
L'homme y passe à travers des forêts de symboles  
Qui l'observent avec des regards familiers.  
Comme de longs échos qui de loin se confondent  
Dans une ténébreuse et profonde unité,  
Vaste comme la nuit et comme la clarté,  
Les parfums, les couleurs et les sons se répondent.  
Il est des parfums frais comme des chairs d'enfants,  
Doux comme les hautbois, verts comme les prairies,  
— Et d'autres, corrompus, riches et triomphants,  
Ayant l'expansion des choses infinies,  
Comme l'ambre, le musc, le benjoin et l'encens,  
Qui chantent les transports de l'esprit et des sens.  
— Charles Baudelaire

Matches

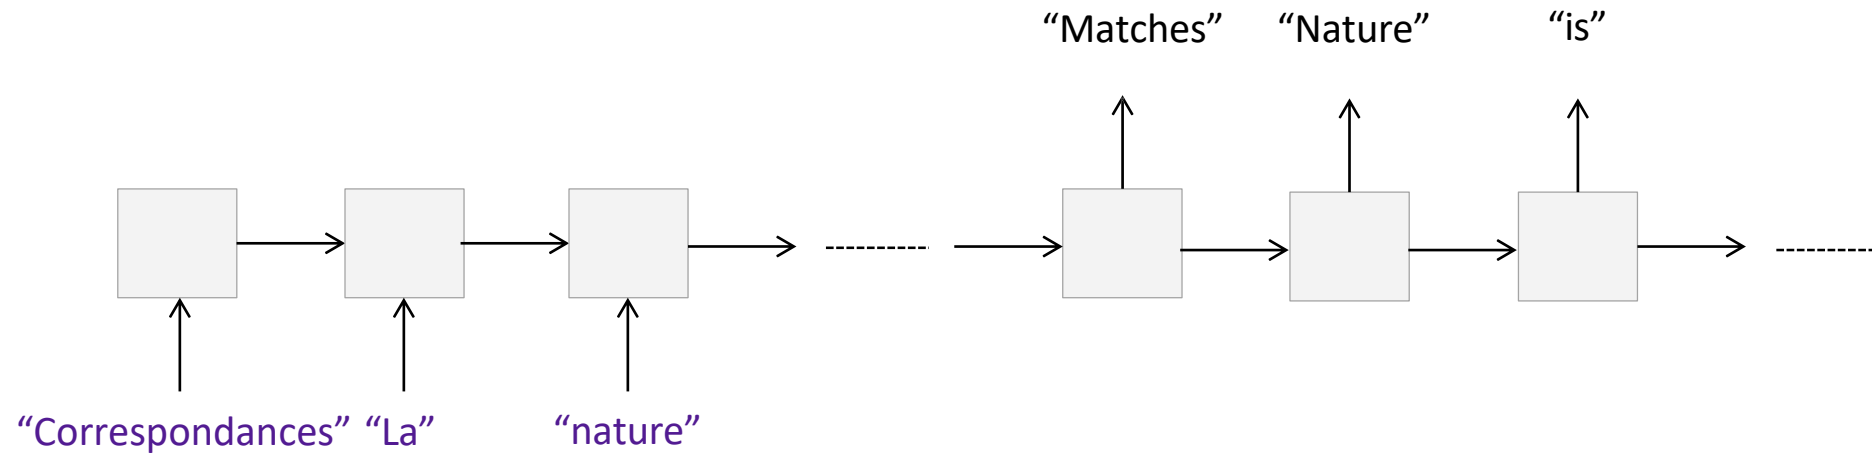
Nature is a temple where living pillars  
Sometimes let out confused words;  
Man goes through symbol forests  
Which observe him with familiar eyes.  
Like long echoes that by far merge  
In a dark and deep unity,  
As vast as the night and as clarity,  
The perfumes, the colors and the sounds answer each  
other.  
There are fresh perfumes like children's flesh,  
Sweet like oboes, green like meadows,  
- And others, corrupt, rich and triumphant,  
Having the expansion of infinite things,  
Like amber, musk, benzoin and incense,  
Who sing the transports of the mind and the senses.  
- Charles Baudelaire

693/5000

<https://translate.google.com/>


# Machine translation

- Multiple input – multiple output (or sequence to sequence)






# Different input output combinations



Movie review – many to one

# Different input output combinations



Many to Many -  
Machine translation  
Dialogue

# Different input output combinations

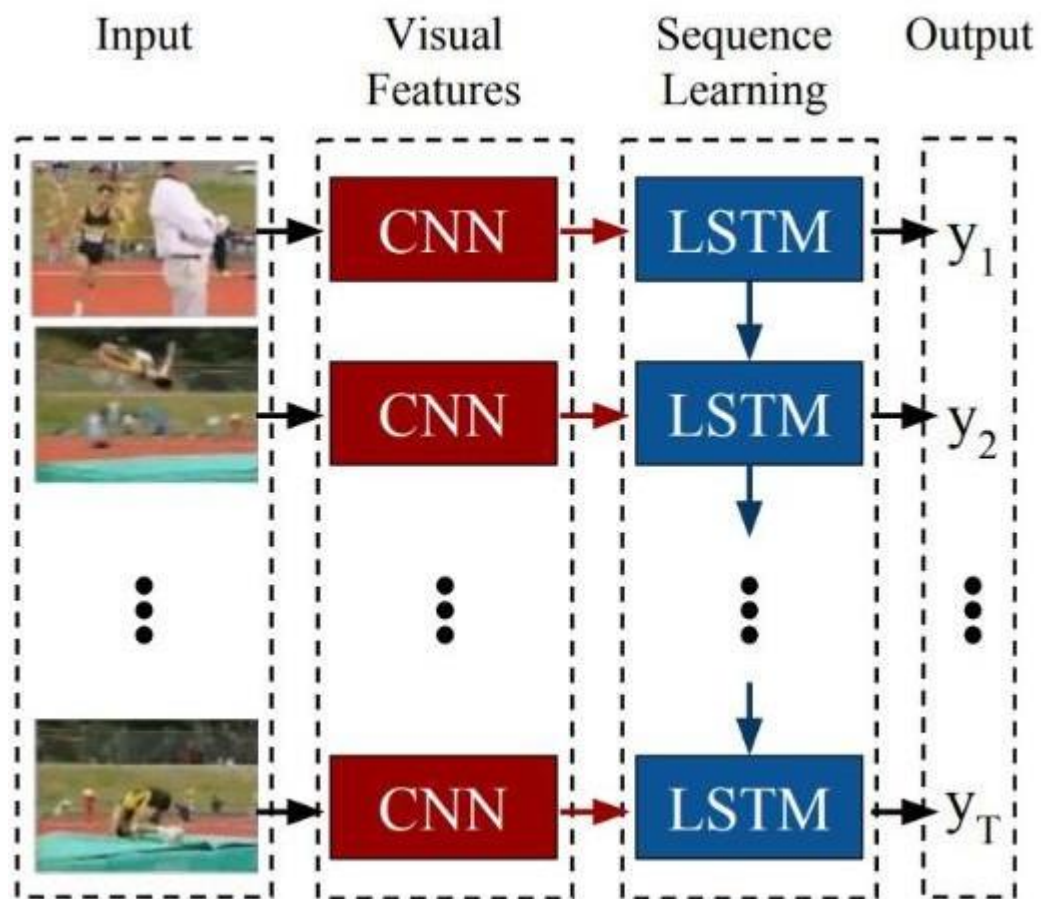
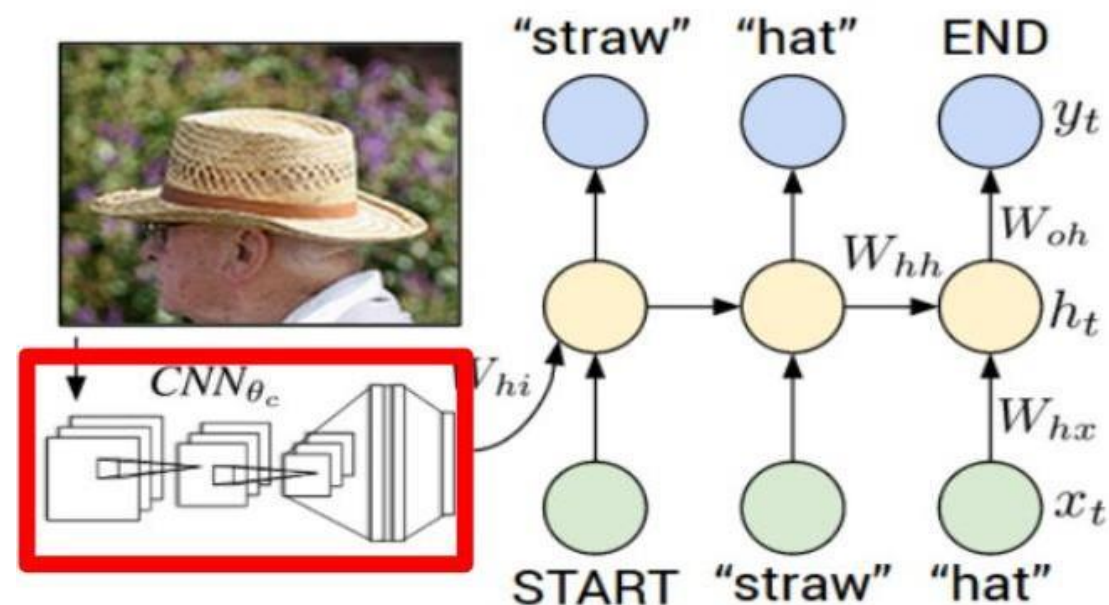


Image and video captioning



Many to many and one to many

Single - Single



Feed-forward Network

Single - Multiple






Image Captioning

Multiple - Single



Sentiment Classification

Multiple - Multiple



Translation

Image Captioning

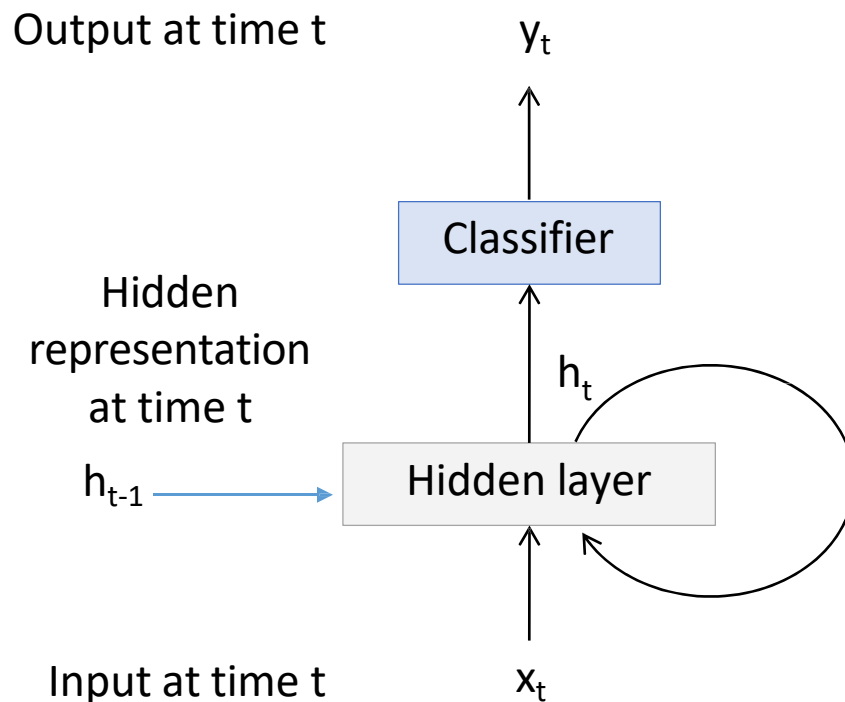
# Sequential learning problem

$$P(w_1, \dots, w_m) = \prod_{i=1}^m P(w_i \mid w_1, \dots, w_{i-1}) \approx \prod_{i=1}^m P(w_i \mid w_{i-(n-1)}, \dots, w_{i-1})$$

The predicted word at a certain point can be inferred given a few previous words

**Markovian Property of languages**

# Idea of hidden state in RNN

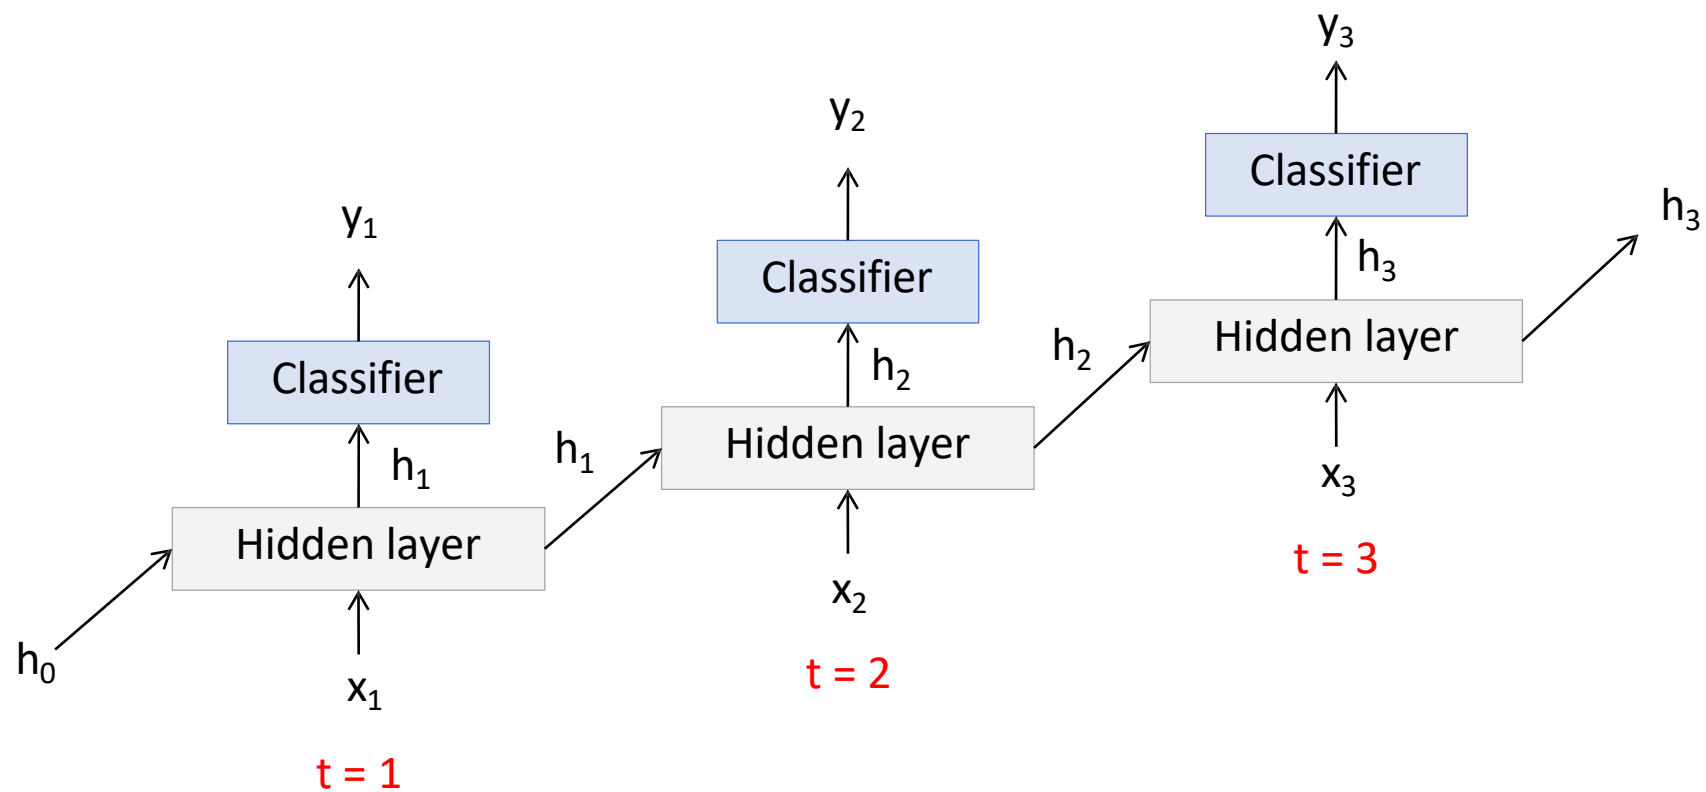


Recurrence:

$$h_t = f_W(x_t, h_{t-1})$$

new state      function of  $W$       input at time  $t$       old state

# Unrolling the RNN



# Formulation

$$h(t) = \sigma (W h(t-1) + U x(t))$$

$$\hat{y}(t) = \text{softmax}(V h(t))$$

$$\hat{P}(x_{t+1} = v_j \mid x_t, \dots, x_1) = \hat{y}_{t,j}$$

Encodes the long-term dependency

$$J^{(t)}(\theta) = - \sum_{j=1}^{|V|} y_{t,j} \log \hat{y}_{t,j}$$

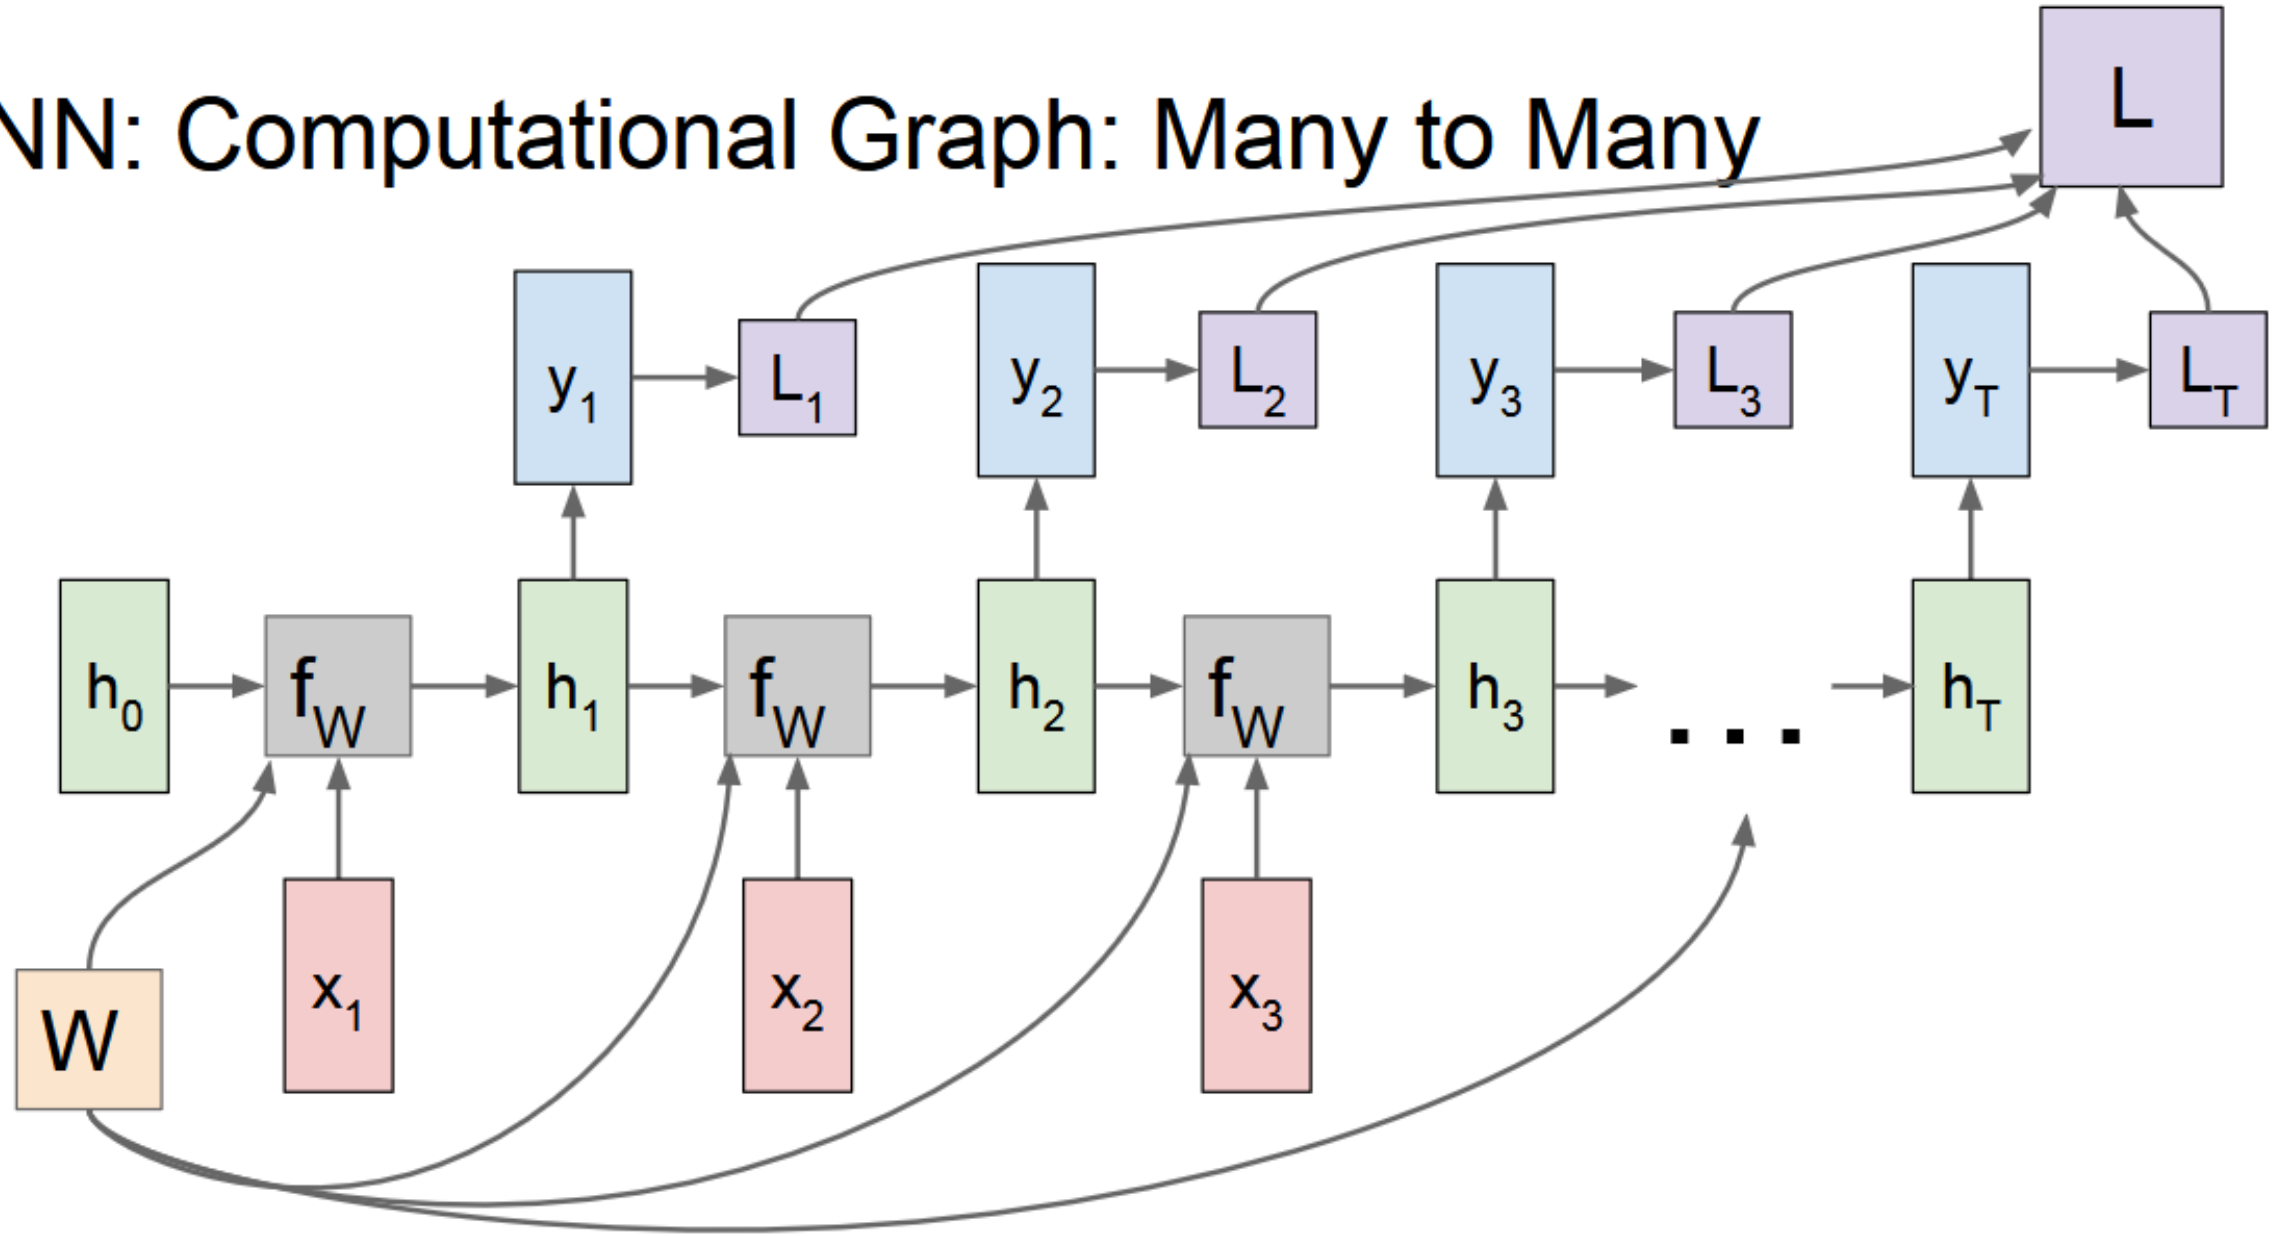
The loss function to be minimized

- ✓ Model for Many to Many prediction
- ✓ One output per time stamp
- ✓  $|V|$  denotes the number of words
- ✓ Training with softmax cross-entropy
- ✓ **The weights are shared for all the time-stamps**

Training is Back-Propagation Through Time



# RNN: Computational Graph: Many to Many



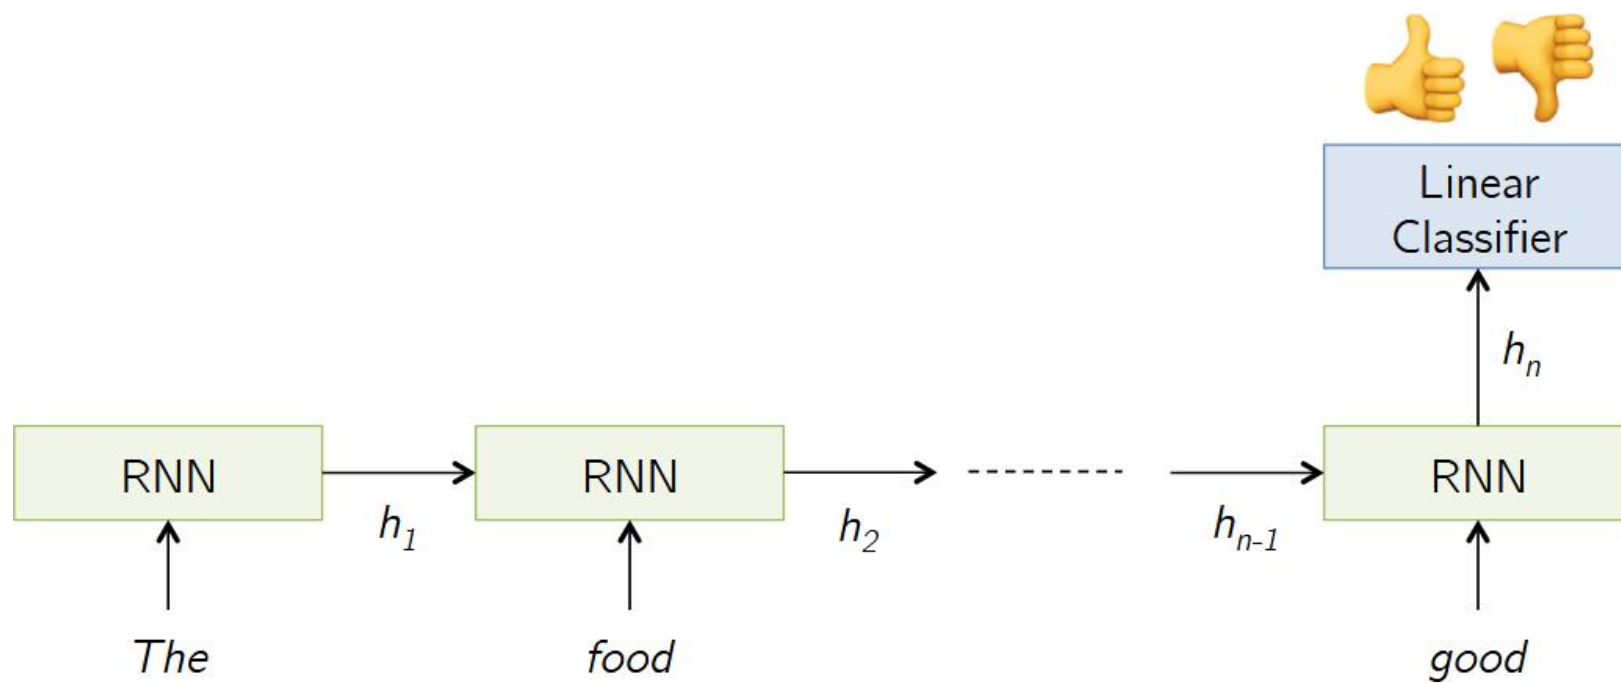
# How to define the classifier

- There could be multiple possibilities for defining the classifier:
  - Ensemble of classifiers for all the states?
  - Single classifier at the final state?
- It is application dependent. Let us see two examples for sentiment classification and captioning

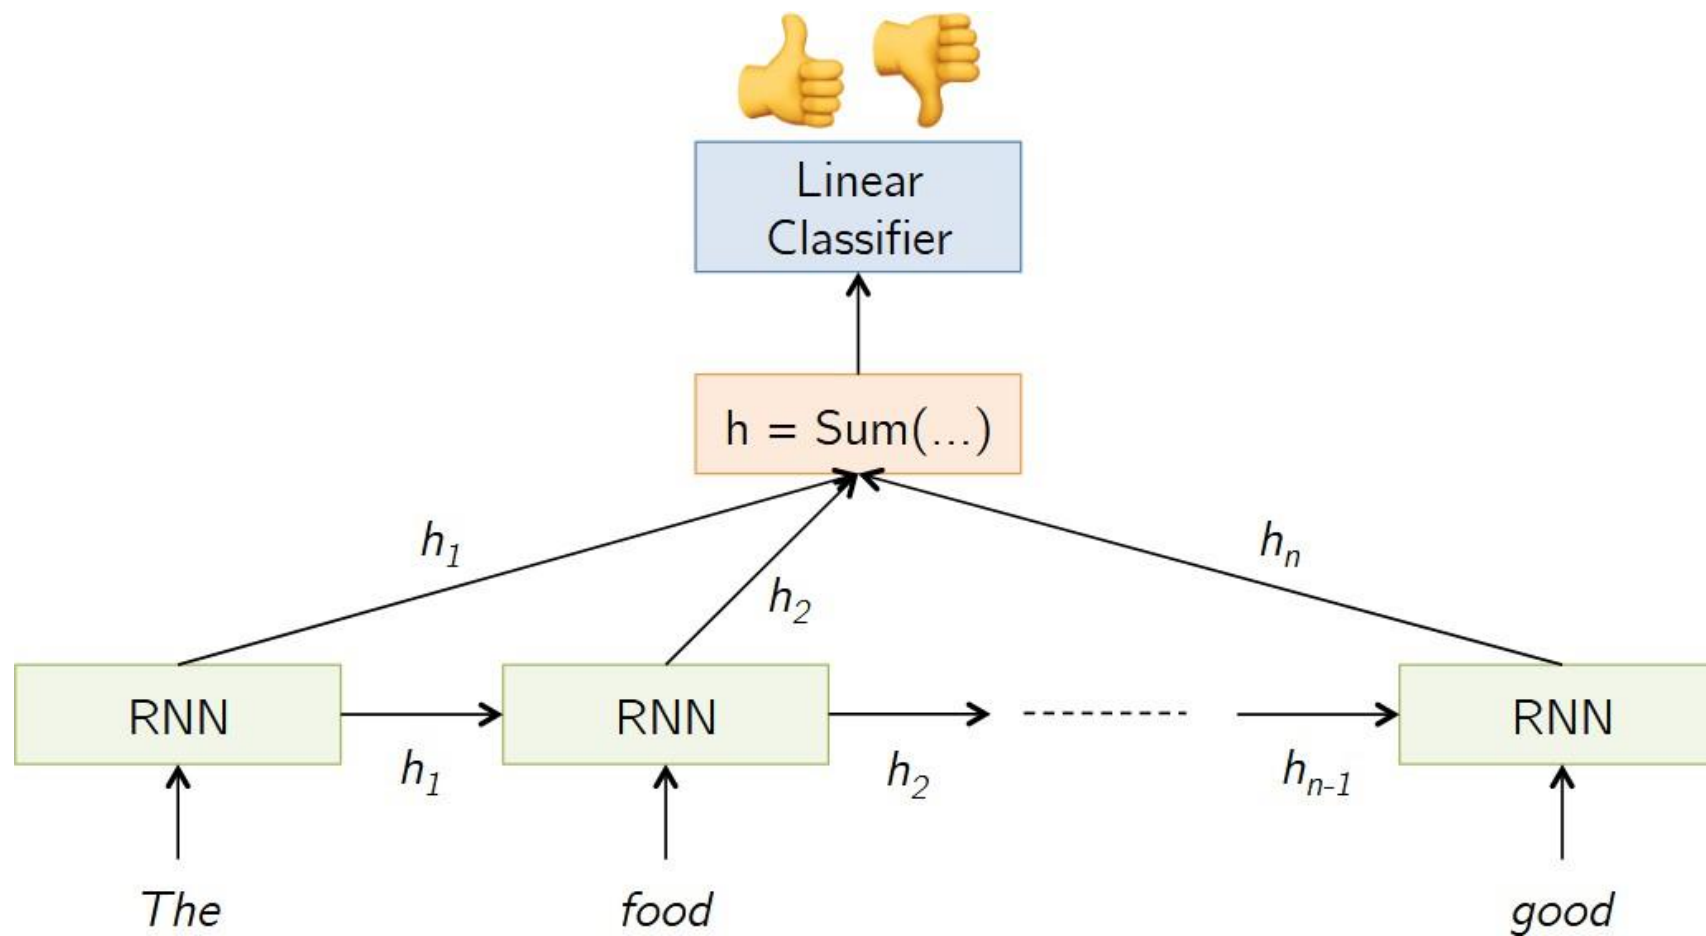
# Example – sentiment classification

- Classify a restaurant review from Yelp! OR movie review from IMDB OR ... as positive or negative
- **Inputs:** Multiple words, one or more sentences
- **Outputs:** Positive / Negative classification
- “The food was really good”
- “The chicken crossed the road because it was uncooked”

# Model-1



# Model-2



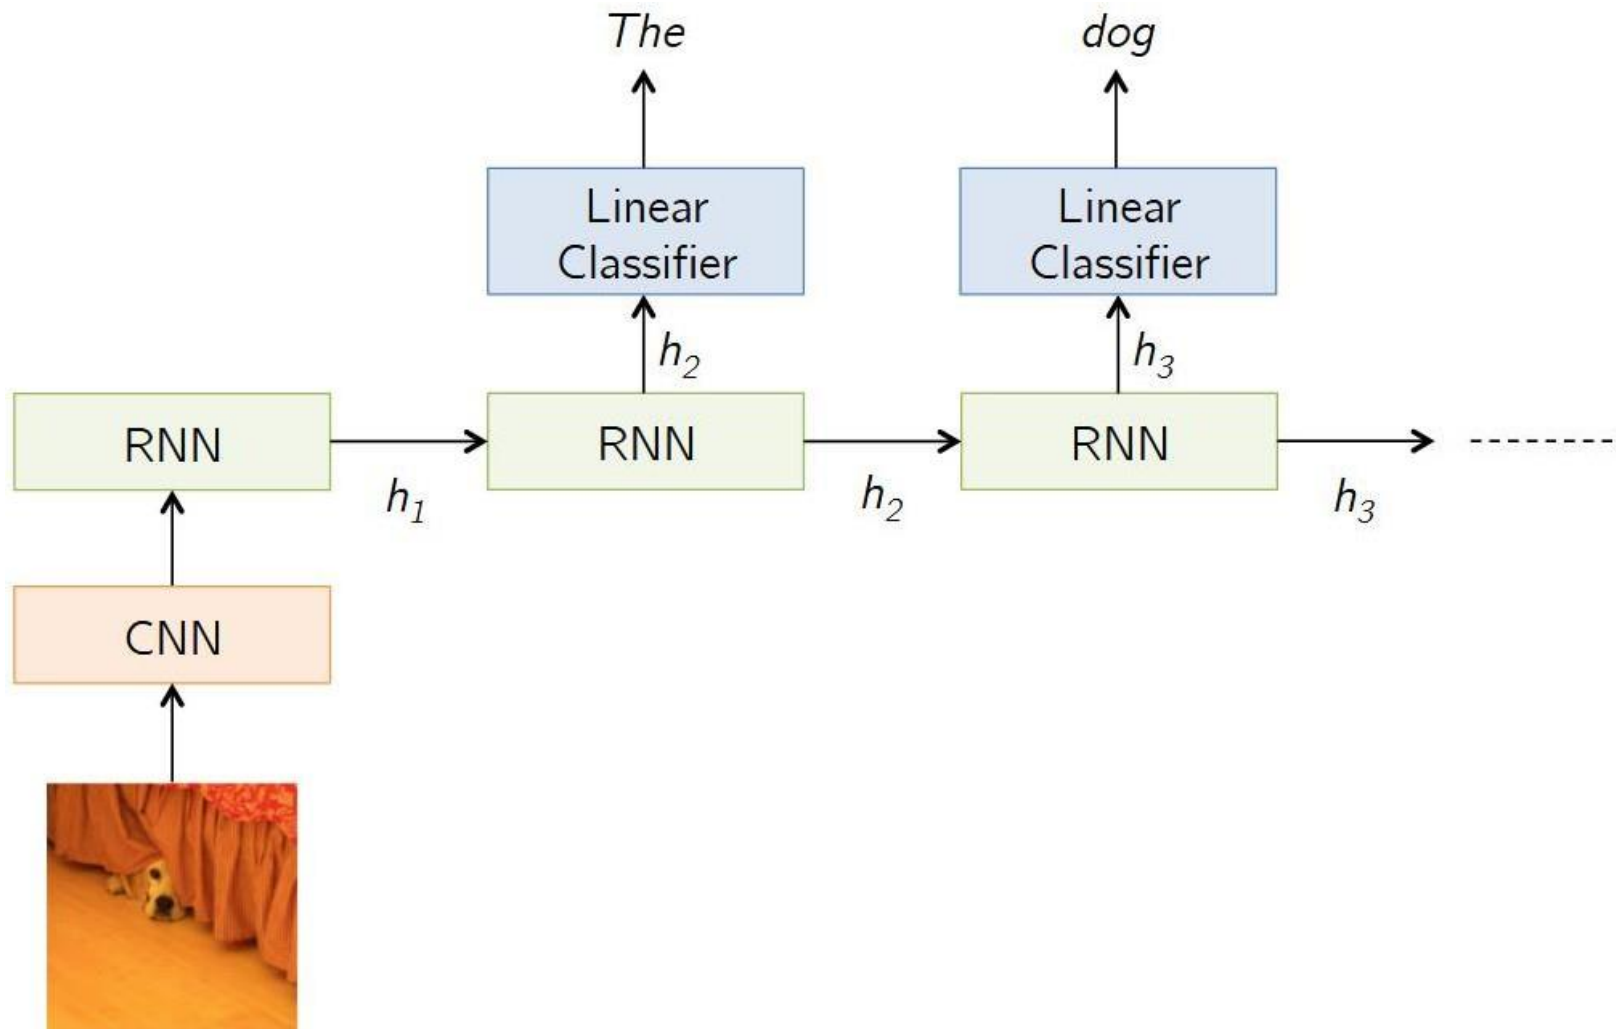
# Image captioning

- Given an image, produce a sentence describing its contents
- **Inputs:** Image feature (from a CNN)
- **Outputs:** Multiple words (let's consider one sentence)



: The dog is hiding

# Model





**A person riding a motorcycle on a dirt road.**



**Two dogs play in the grass.**



**A herd of elephants walking across a dry grass field.**



**A group of young people playing a game of frisbee.**



**Two hockey players are fighting over the puck.**



**A close up of a cat laying on a couch.**







# Some highlights

- RNN takes the previous output or hidden state as inputs at every time together with the usual input. This means that the composite input at time  $t$  has some historical information about the happenings at time  $T < t$ .
- RNNs gained popularity in sequence modeling without forgetting important past information as their intermediate values (state) can store information about the past inputs for a time that is not fixed a priori.

# Back propagation through time

- Considered to be the standard method used to train RNNs.
- The unfolded network (used during forward pass) is treated as one feed-forward network that accepts the whole time series as input.
- The weight updates are computed for each copy in the unfolded network, then accumulated and applied to the model weights and biases.

# Forward and backward pass



$$\begin{aligned}h_t &= f(x_t, h_{t-1}, w_h), \\o_t &= g(h_t, w_o),\end{aligned}$$

where  $f$  and  $g$  are transformations of the hidden layer and the output layer, respectively. Hence, we have a chain of values  $\{\dots, (x_{t-1}, h_{t-1}, o_{t-1}), (x_t, h_t, o_t), \dots\}$  that depend on each other via recurrent computation. The forward propagation is fairly straightforward. All we need is to loop through the  $(x_t, h_t, o_t)$  triples one time step at a time. The discrepancy between output  $o_t$  and the desired target  $y_t$  is then evaluated by an objective function across all the  $T$  time steps as

$$L(x_1, \dots, x_T, y_1, \dots, y_T, w_h, w_o) = \frac{1}{T} \sum_{t=1}^T l(y_t, o_t).$$

$$\begin{aligned}\frac{\partial L}{\partial w_h} &= \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial w_h} \\&= \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial o_t} \frac{\partial g(h_t, w_o)}{\partial h_t} \frac{\partial h_t}{\partial w_h}.\end{aligned}$$

$$\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_h}.$$

To derive the above gradient, assume that we have three sequences  $\{a_t\}$ ,  $\{b_t\}$ ,  $\{c_t\}$  satisfying  $a_0 = 0$  and  $a_t = b_t + c_t a_{t-1}$  for  $t = 1, 2, \dots$ . Then for  $t \geq 1$ , it is easy to show

$$a_t = b_t + \sum_{i=1}^{t-1} \left( \prod_{j=i+1}^t c_j \right) b_i.$$

By substituting  $a_t$ ,  $b_t$ , and  $c_t$  according to

$$\begin{aligned} a_t &= \frac{\partial h_t}{\partial w_h}, \\ b_t &= \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h}, \\ c_t &= \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial h_{t-1}}, \end{aligned}$$

$$\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \sum_{i=1}^{t-1} \left( \prod_{j=i+1}^t \frac{\partial f(x_j, h_{j-1}, w_h)}{\partial h_{j-1}} \right) \frac{\partial f(x_i, h_{i-1}, w_h)}{\partial w_h}.$$

The repeated multiplication of the gradients cause the derivative to explode

# The exploding gradient problem

Imagine an RNN that just does:

$$h_1 = w \cdot h_0,$$

$$h_2 = w \cdot h_1,$$

$$h_3 = w \cdot h_2,$$

...

where:

- $h_t$  is the hidden state at timestep  $t$  (just a single number here).
- $w$  is a single learned parameter (instead of a matrix).
- $h_0$  is some initial state.

We define a **loss**  $L$  that depends on the final hidden state (say,  $h_3$  in this tiny example). For simplicity, let's say  $L$  is just  $L(h_3)$ . When we do Backpropagation Through Time, we need  $\frac{\partial L}{\partial w}$ .

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial h_3} \times \frac{\partial h_3}{\partial w}.$$

$$1. \ h_3 = w \cdot h_2$$

$$\frac{\partial h_3}{\partial w} = h_2 + w \frac{\partial h_2}{\partial w} \quad (\text{Product Rule})$$

$$2. \ h_2 = w \cdot h_1$$

$$\frac{\partial h_2}{\partial w} = h_1 + w \frac{\partial h_1}{\partial w}.$$

$$3. \ h_1 = w \cdot h_0$$

$$\frac{\partial h_1}{\partial w} = h_0 \quad (\text{no extra term, since } h_0 \text{ is constant}).$$

Putting these together, we go from the inside out:

$$\frac{\partial h_2}{\partial w} = h_1 + w \underbrace{\frac{\partial h_1}{\partial w}}_{=h_0} = w \cdot h_0 + w \cdot h_0 = 2 w h_0.$$


Then,

$$\frac{\partial h_3}{\partial w} = h_2 + w \frac{\partial h_2}{\partial w} = \underbrace{w \cdot h_1}_{=w^2 h_0} + w (2 w h_0) = w^2 h_0 + 2 w^2 h_0 = 3 w^2 h_0.$$

$$\frac{\partial L}{\partial w} = \delta_3 \times (3 w^2 h_0).$$

For **3 timesteps**, we got a factor of **3** and a factor of  $w^2$ . If we had  $T$  **timesteps**, you'd see terms that grow roughly like  $T \cdot w^{T-1}$ . If  $|w| > 1$ , this can become **very large** for even moderate  $T$ .


# Solutions



Instead of processing the entire sequence, process the forward/backward passes in chunks




# Solutions




# Short term dependency

Language model trying to predict the next word based on the previous ones



# Long term dependency

I grew up in India... I speak fluent Hindi.



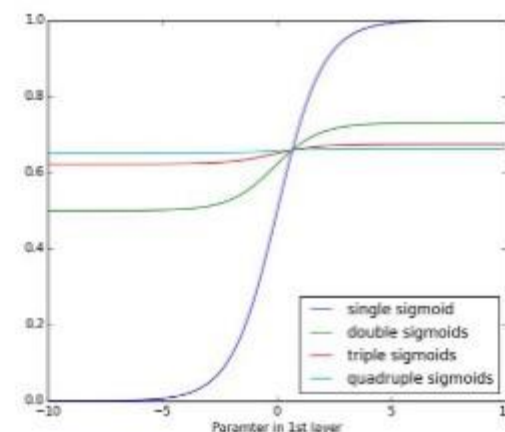

RNN fails to preserve the long term dependency due to the vanishing gradient problem

We should distinguish between what to remember and what to forget more wisely

LSTM (Long short term memory) to the rescue



# The RNN problem revisited

- Problem: can't capture long-term dependencies due to vanishing/exploding gradients during backpropagation



$$\begin{aligned} h^{(t)} &= \sigma(w_c \cdot c^{(t)}) \\ c^{(t)} &= \sigma(w_r \cdot c^{(t-1)} + w_x \cdot x^{(t)}) \end{aligned}$$

$$\begin{aligned} h^{(3)} &= \sigma(w_c \cdot c^{(3)}) \\ &= \sigma(w_c \cdot \sigma(w_x \cdot x^{(3)} + w_r \cdot c^{(2)})) \\ &= \sigma(w_c \cdot \sigma(w_x \cdot x^{(3)} + w_r \cdot \sigma(w_x \cdot x^{(2)} + w_r \cdot c^{(1)})))))) \\ &= \sigma(w_c \cdot \sigma(w_x \cdot x^{(3)} + w_r \cdot \sigma(w_x \cdot x^{(2)} + w_r \cdot \sigma(w_x \cdot x^{(1)} + w_r \cdot c^{(0)})))))) \end{aligned}$$



-  Tanh function
-   $h_t$  new hidden state
-   $h_{t-1}$  previous hidden state
-   $x_t$  input
-  concatenation

How RNN works! – trying to process all the words and update memory on that basis

# Some intuition

Customers Review 2,491




Thanos

September 2018

Verified Purchase

**Amazing! This box of cereal gave me a perfectly balanced breakfast, as all things should be. I only ate half of it but will definitely be buying again!**




A Box of Cereal  
**\$3.99**

We look for certain words, while ignoring most of the other words


# LSTM

**Central Idea:** A **memory cell** (interchangeably **block**) which can maintain its state over time, consisting of an explicit memory (aka the **cell state vector**) and **gating units** which regulate the information flow into and out of the memory.



LSTM Memory Cell



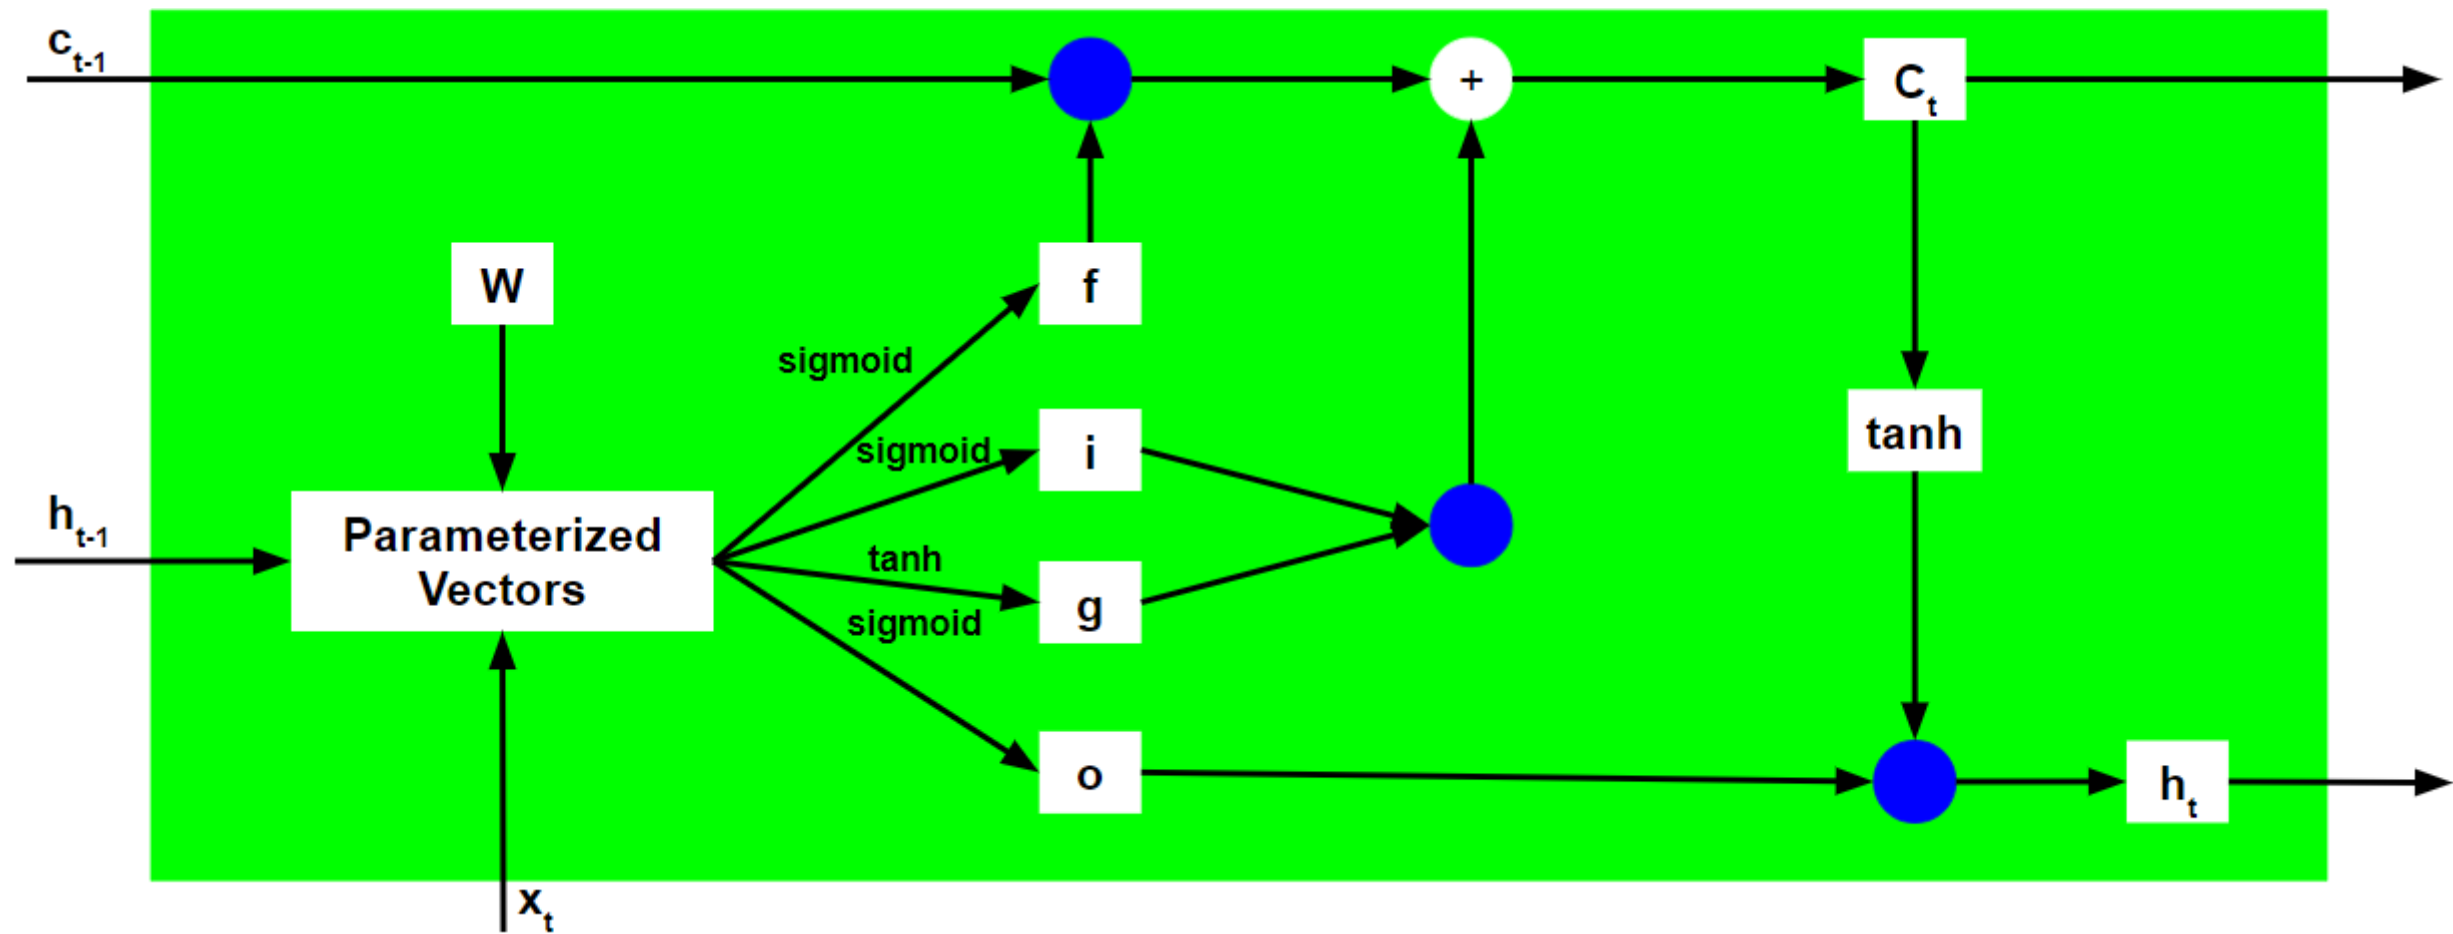


## 1. Hidden State ( $h_t$ )

- **Short-Term Signal:** The hidden state is effectively the “working” or “immediate” memory. It carries information that is directly used to predict the output at the current timestep (and to help compute the next states).
- **Frequently Updated:** Because it directly connects to the network’s output,  $h_t$  changes more rapidly and captures shorter-term patterns.

## 2. Cell State ( $c_t$ )

- **Long-Term Memory:** The cell state is designed to propagate information over much longer sequences with fewer modifications. It flows more linearly through the network (with some gating).
- **Reduced Vanishing/Exploding Gradients:** By providing a nearly linear path for gradients,  $c_t$  helps mitigate the vanishing and exploding gradient problems, which commonly arise in vanilla RNNs.



## Forget Gate

**Purpose:** Controls which information from  $c_{t-1}$  should be discarded before forming  $c_t$ .

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

- $\sigma(\cdot)$  is the sigmoid function, which outputs values between 0 and 1.
- $f_t$  is multiplied elementwise with the previous cell state  $c_{t-1}$ . Values close to 0 mean "forget most of this," while values close to 1 mean "keep most of this."

**Applied to  $c_{t-1}$ :**

$$\tilde{c}_{t-1} = f_t \odot c_{t-1}$$

# Input Gate

**Purpose:** Determines how much new information from the current timestep's input (and the previous hidden state) is added to the cell state.

This gate is actually split into two parts—an **input gate** and a **candidate cell state**. We first compute a “candidate” update, then use the input gate to decide how much of this candidate actually goes into  $c_t$ .

1. **Candidate Cell State** ( $\tilde{c}_t$ ):

$$\tilde{c}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

This is the new, potential information that could be added to the cell state.

2. **Input Gate** ( $i_t$ ):

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

Outputs a value between 0 and 1 for each component, controlling how much of  $\tilde{c}_t$  is added.

**Combined effect:**

$$i_t \odot \tilde{c}_t$$

This value is added to the (partially) “forgotten” previous cell state  $\tilde{c}_{t-1}$  to form the new cell state.

With the forget gate and the input gate, we can update  $c_t$  as follows:

$$c_t = \underbrace{(f_t \odot c_{t-1})}_{\text{forgotten old state}} + \underbrace{(i_t \odot \tilde{c}_t)}_{\text{new information}}$$

## Output Gate

**Purpose:** Determines how much of the new cell state  $c_t$  will be used to compute the hidden state  $h_t$ .

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

- $\sigma(\cdot)$  is the sigmoid function (again producing values between 0 and 1).

We then apply a **tanh** function to the newly updated cell state  $c_t$  to push values between -1 and +1, and multiply elementwise by  $o_t$  to get the final hidden state  $h_t$ :

$$h_t = o_t \odot \tanh(c_t)$$

# Summarizing the equations

$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma (W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$




# Training LSTM


- Backpropagation through time
- Each cell has many parameters ( $W_f$ ,  $W_i$ ,  $W_C$ ,  $W_o$ )
  - Generally requires lots of training data.
  - Requires lots of compute time that exploits GPU clusters.

# Deep LSTM

- Deep LSTMs can be created by stacking multiple LSTM layers vertically, with the output sequence of one layer forming the input sequence of the next (in addition to recurrent connections within the same layer)
- Increases the number of parameters - but given sufficient data, performs significantly better than single-layer LSTMs (Graves et al. 2013)
- Dropout usually applied only to non-recurrent edges, including between layers



# Bi-directional recurrent model



1. **Forward Pass:**

- Processes the input sequence  $(x_1, x_2, \dots, x_T)$  from  $t = 1$  to  $t = T$ .
- Produces a set of "forward" hidden states  $(\vec{h}_1, \vec{h}_2, \dots, \vec{h}_T)$ .

2. **Backward Pass:**

- Processes the sequence in reverse  $(x_T, x_{T-1}, \dots, x_1)$ .
- Produces a set of "backward" hidden states  $(\overleftarrow{h}_T, \overleftarrow{h}_{T-1}, \dots, \overleftarrow{h}_1)$ .


3. **Combining Forward and Backward States:**

At each timestep  $t$ , the forward hidden state  $\vec{h}_t$  and the backward hidden state  $\overleftarrow{h}_t$  are combined (often by **concatenation** or **addition**) to form a single output representation:

$$h_t = \left[ \vec{h}_t; \overleftarrow{h}_t \right]$$

where  $;$  denotes concatenation. This final  $h_t$  captures information from **both past and future context** relative to timestep  $t$ .

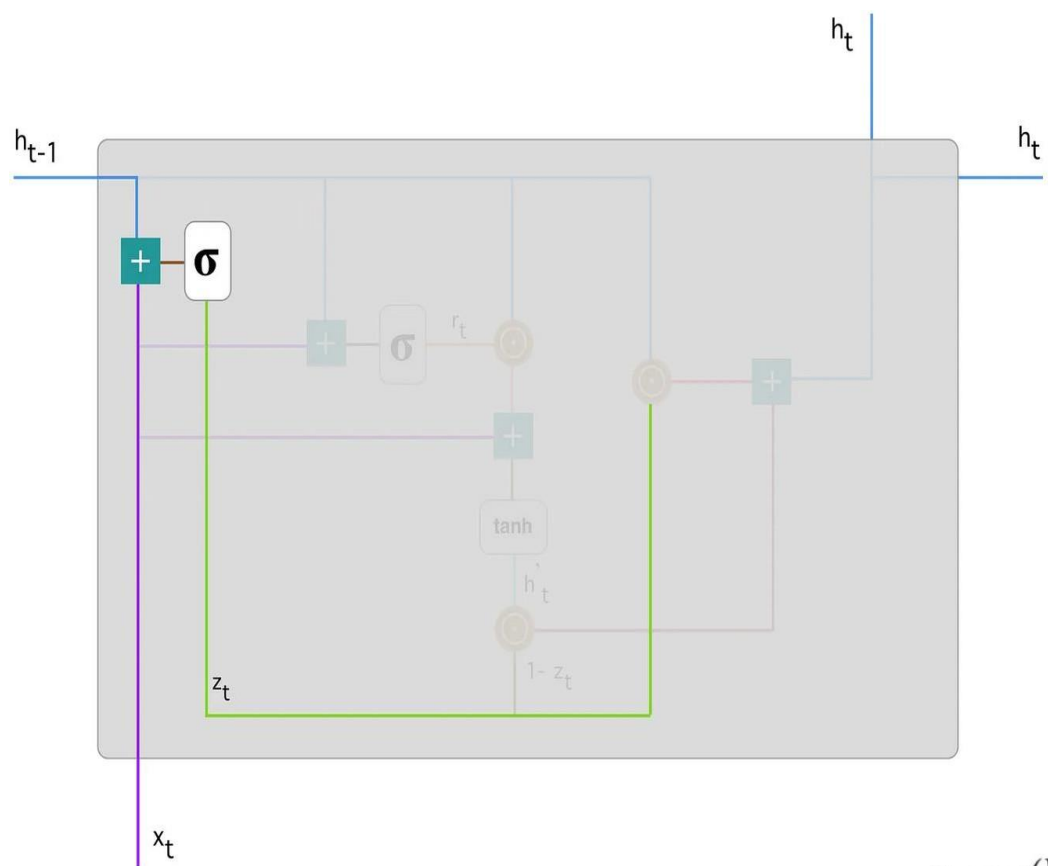
# Gated recurrent units



GRU's got rid of the cell state and used the hidden state to transfer information. It also only has two gates, a reset gate and update gate.

# Update gate

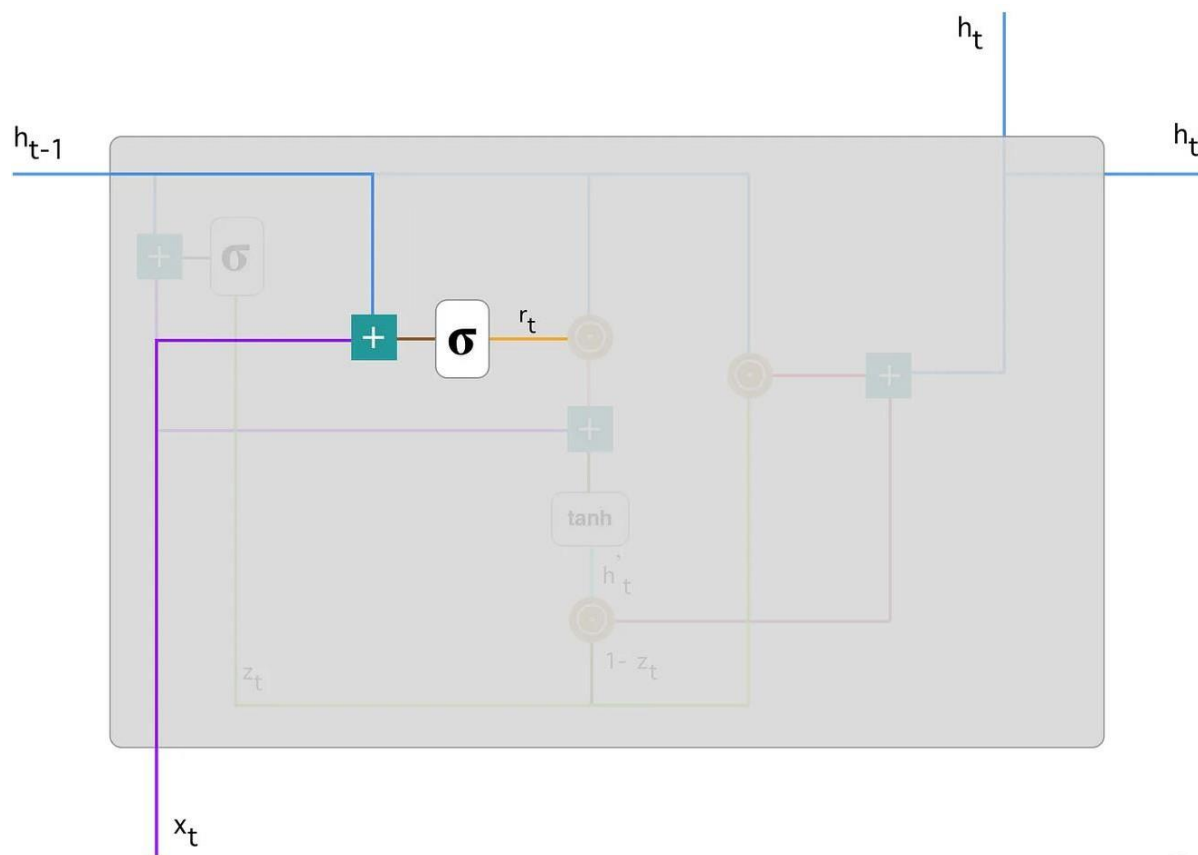
The update gate decides how much of the previous hidden state  $h_{t-1}$  should be carried over to the new hidden state. A higher value of  $z_t$  means “keep more of the past information,” while a lower value means “let in more of the new candidate information.”



$$z_t = \sigma(W^{(z)}x_t + U^{(z)}h_{t-1})$$

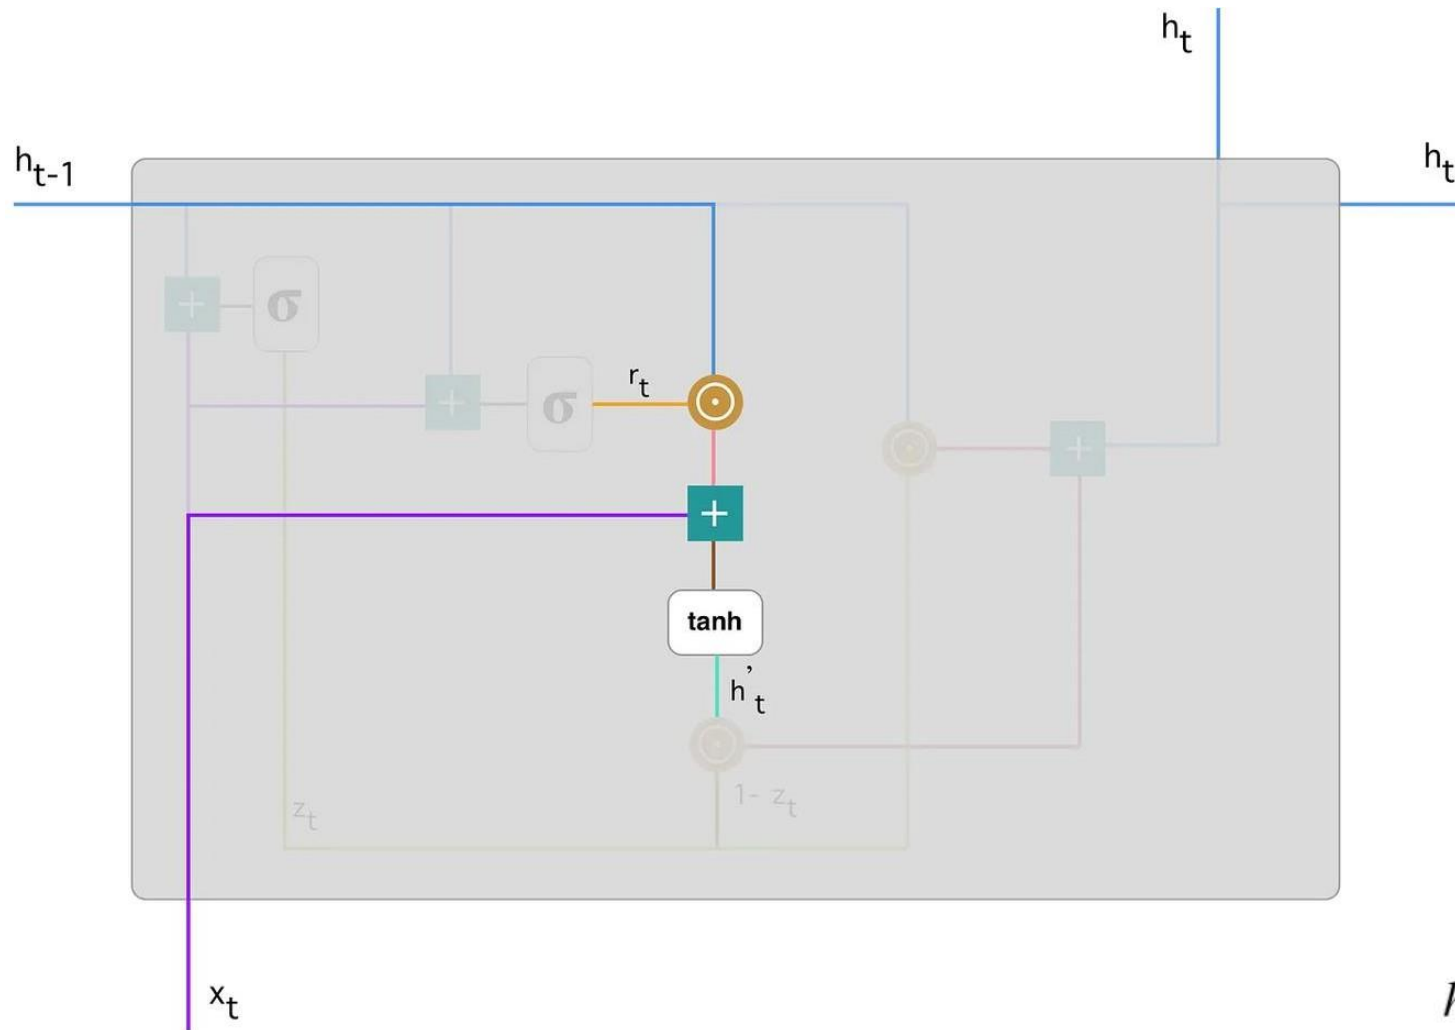
# Reset gate

The reset gate determines how much of the past hidden state  $h_{t-1}$  to “forget” (or reset) before calculating the new candidate hidden state. A lower value of  $r_t$  means the model puts less emphasis on the previous hidden state, effectively forgetting more past context.



$$r_t = \sigma(W^{(r)}x_t + U^{(r)}h_{t-1})$$

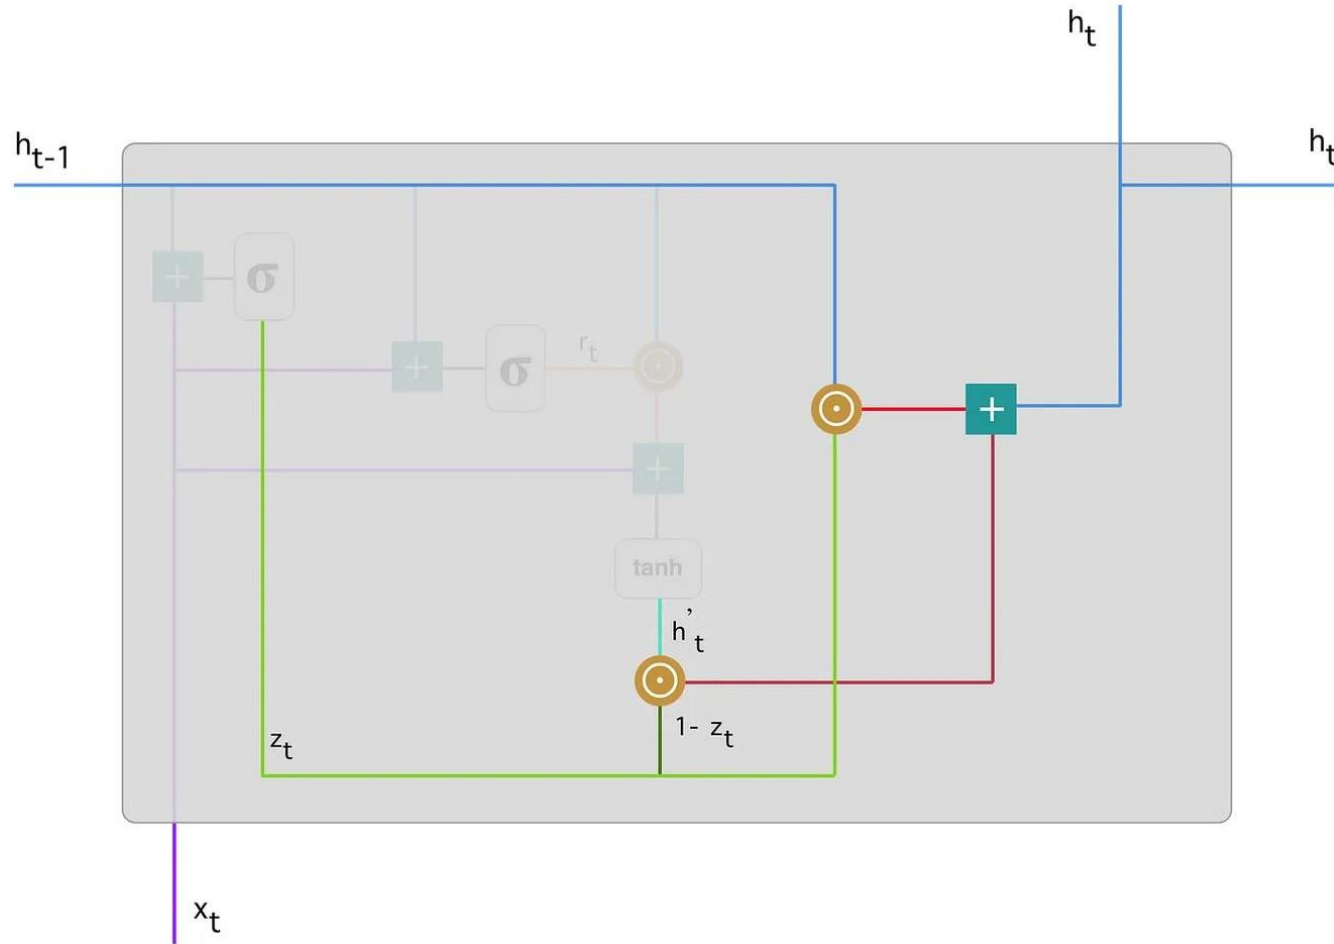
# Creating intermediate memory



$$h'_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$



# Outputting the next hidden state



$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot h'_t$$

```
initialize hidden state  $h_0 = 0$   
for t in 1 to T:  
     $z_t = \text{sigmoid}(W_z * [h_{t-1}, x_t] + b_z)$   
     $r_t = \text{sigmoid}(W_r * [h_{t-1}, x_t] + b_r)$   
     $h_{t\_candidate} = \tanh(W_h * [r_t * h_{t-1}, x_t] + b_h)$   
     $h_t = (1 - z_t) * h_{t-1} + z_t * h_{t\_candidate}$ 
```

Consider a linear autoencoder with:

- An input layer of dimension  $d$ .
- A hidden/"bottleneck" layer of dimension  $k$  (with  $k < d$ ).
- A reconstruction layer also of dimension  $d$ .

Let the encoder be described by a weight matrix  $W \in \mathbb{R}^{k \times d}$  and bias  $\mathbf{b} \in \mathbb{R}^k$ , and the decoder by another weight matrix  $V \in \mathbb{R}^{d \times k}$  and bias  $\mathbf{c} \in \mathbb{R}^d$ . For a dataset  $\{\mathbf{x}^{(n)}\}_{n=1}^N$  with  $\mathbf{x}^{(n)} \in \mathbb{R}^d$ , define the (mean-squared) reconstruction loss as:

$$\mathcal{L} = \sum_{n=1}^N \|\mathbf{x}^{(n)} - \hat{\mathbf{x}}^{(n)}\|^2,$$

where

$$\hat{\mathbf{x}}^{(n)} = V (W \mathbf{x}^{(n)} + \mathbf{b}) + \mathbf{c}.$$

1. **Show** that, after removing any mean offset of the data via preprocessing (so you can assume  $\sum_{n=1}^N \mathbf{x}^{(n)} = \mathbf{0}$  and drop the bias terms  $\mathbf{b}, \mathbf{c}$  for simplicity), the weight matrices  $W$  and  $V$  that **minimize**  $\mathcal{L}$  satisfy

$$V = W^\top.$$

2. **Prove** that the optimal matrix  $W^*$  (which also implies  $V^* = (W^*)^\top$ ) forms an orthonormal basis for the **top**  $k$  principal components of the data covariance matrix

$$\Sigma = \frac{1}{N} \sum_{n=1}^N \mathbf{x}^{(n)} (\mathbf{x}^{(n)})^\top.$$

In other words, the rows of  $W^*$  span the same subspace as the eigenvectors of  $\Sigma$  associated with its largest  $k$  eigenvalues.